



Debugging Unikernel Operating Systems

Kareem Ahmad, Alan Dearle, Jon Lewis, Ward Jaradat

School of Computer Science

University of St Andrews

Email: kareemahmad@protonmail.com



Overview

- In this talk we report on an undergraduate led project to develop debugger for unikernels running on Xen.
- Unikernels are challenging to debug as there are not many production ready debuggers for unikernels.
- Specifically, we focused on debugging support for the [Stardust](#) unikernel.
- This work is applicable to any unikernel written in C and hosted on Xen.

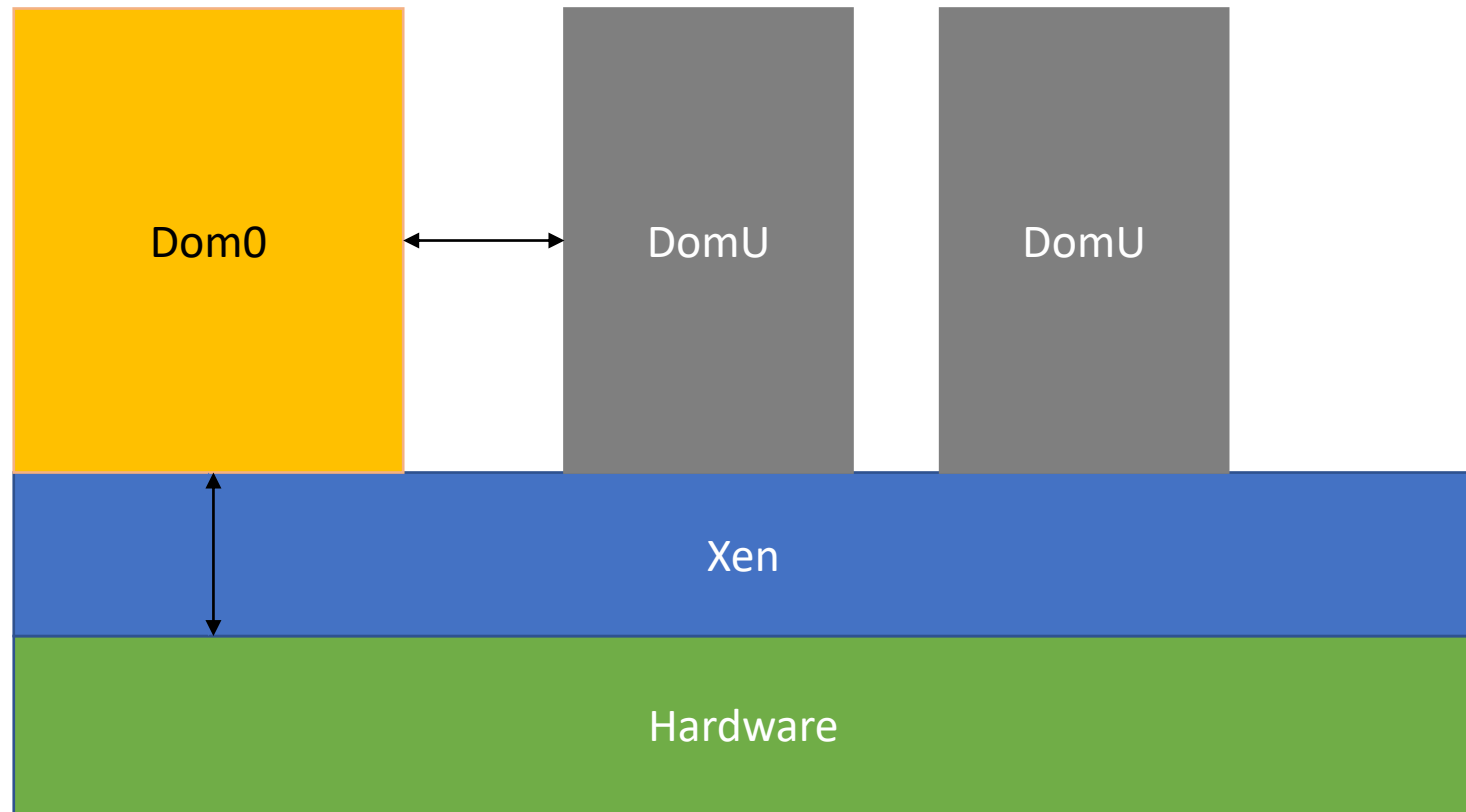


What Is A Unikernel?

- Specialised Operating System that runs directly above the Hypervisor (such as Xen)
- Single image that contains the OS, an application plus any required libraries
- Small image size (400KB including application)
- Fast to boot
- Fast to deploy
- Fast to provision



What is Xen?





The Problem

- Unikernels are difficult to debug:
 - The kernel and the application are compiled into a single image requiring embedded support.
 - An independent debugging context is needed in order to provide isolation and the ability to stop and start the Operating System.
 - Unikernels may not be designed for compatibility with conventional debugging tools like gdb
 - Gdb commonly makes use of Unix process structures, ptrace, and library calls



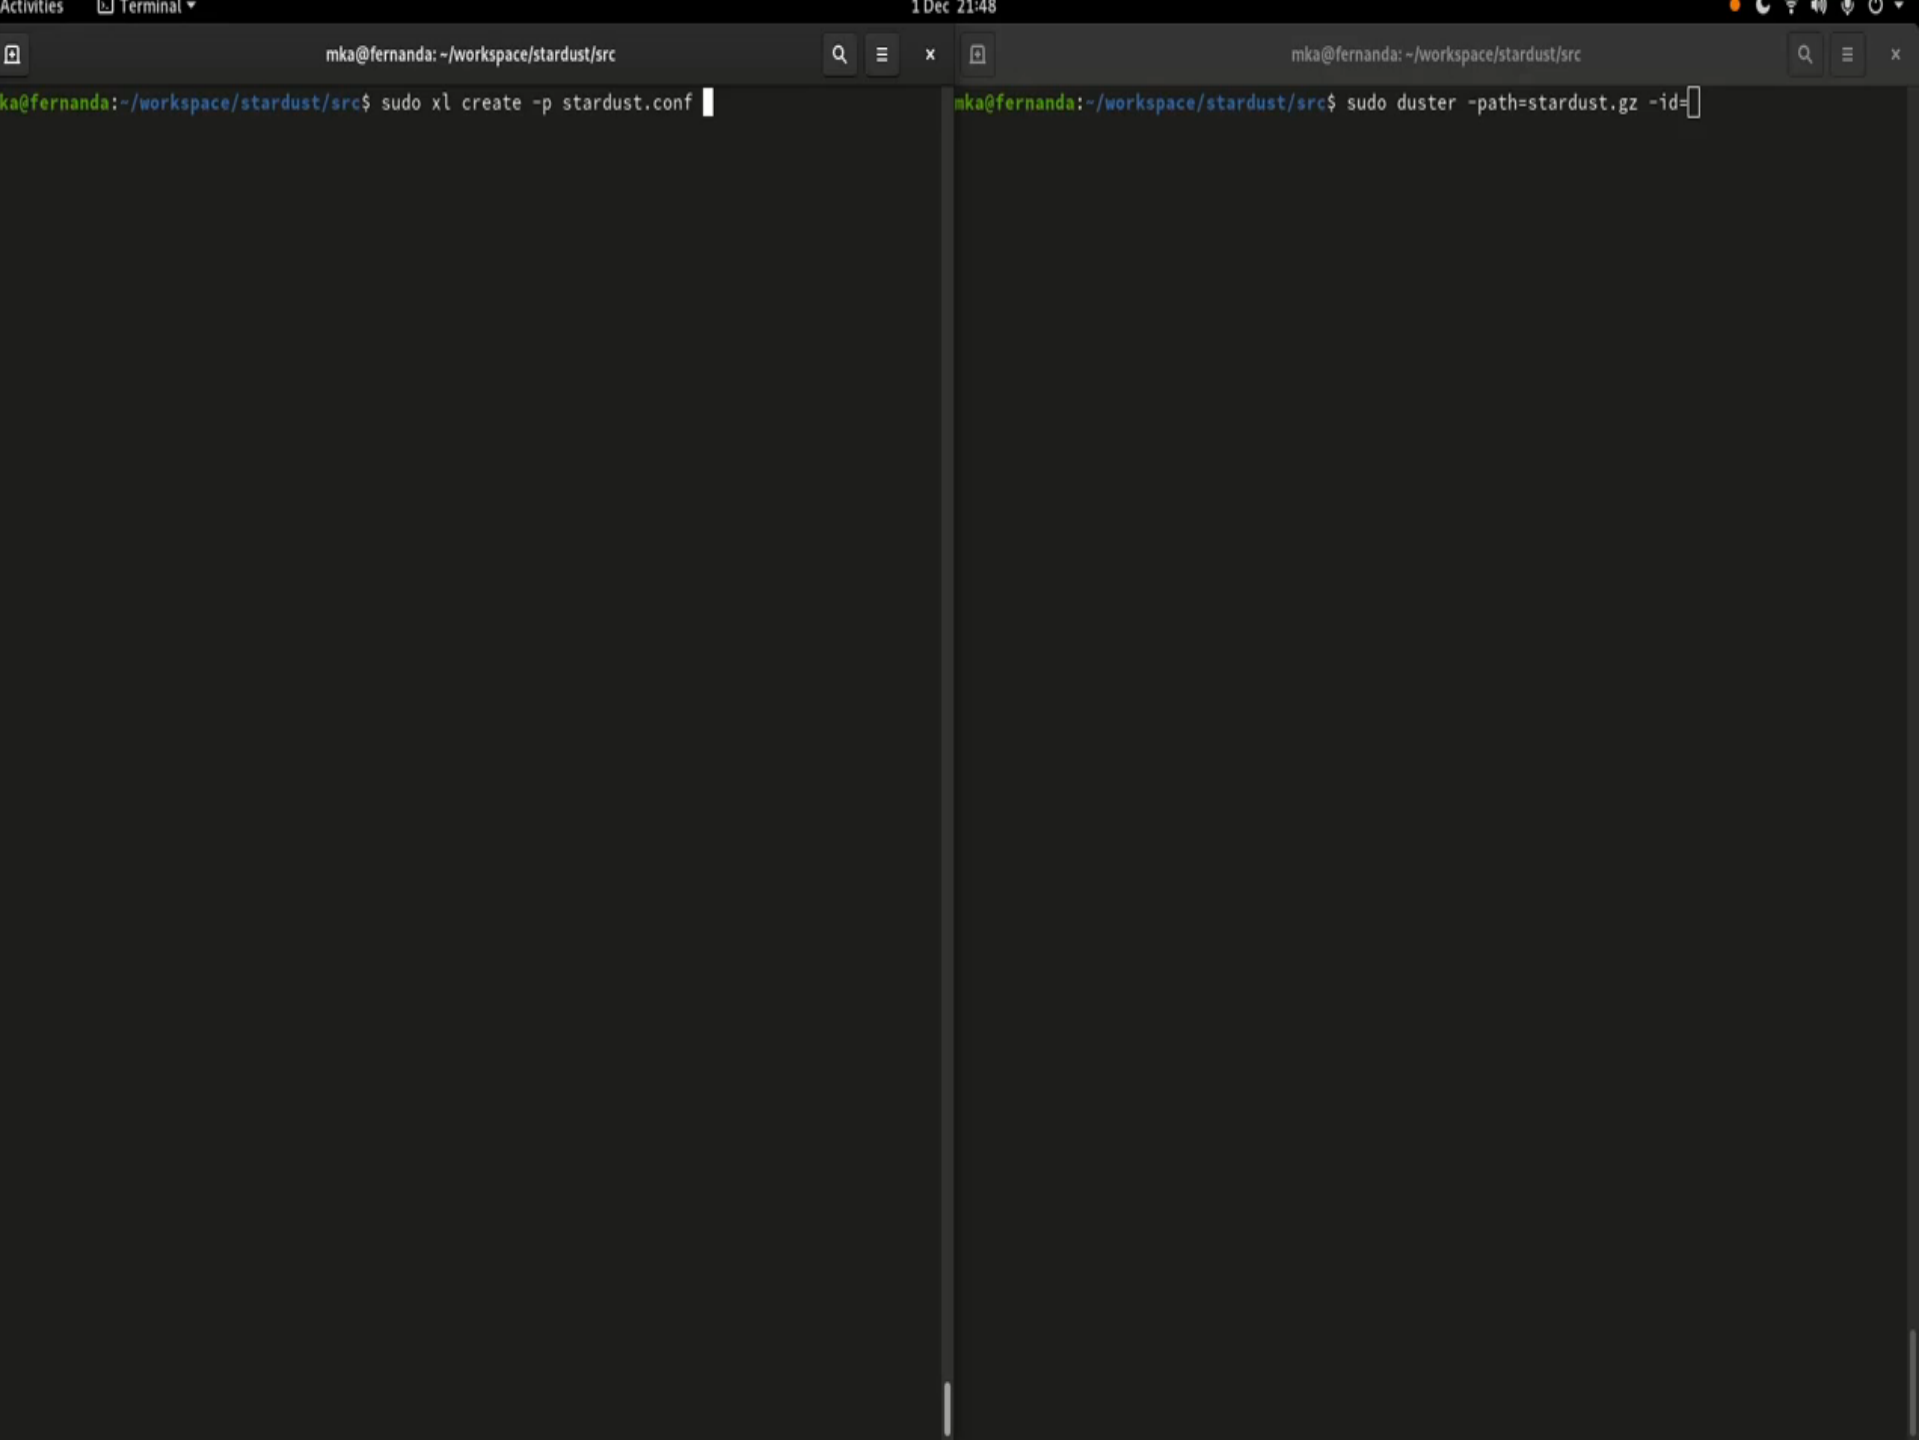

Approach

- The approach taken was based on xendbg developed at Nccgroup by Michael Spencer
- Algorithms and techniques used by Duster and xendbg are very similar.
- Duster focused on allowing the developer to use source-level constructs during debugging
- Uses Xen's Virtual Machine Introspection API for interacting with the Unikernel
- Written in C and Go
- Uses standard DWARF file format to get symbolic information (e.g. variables)



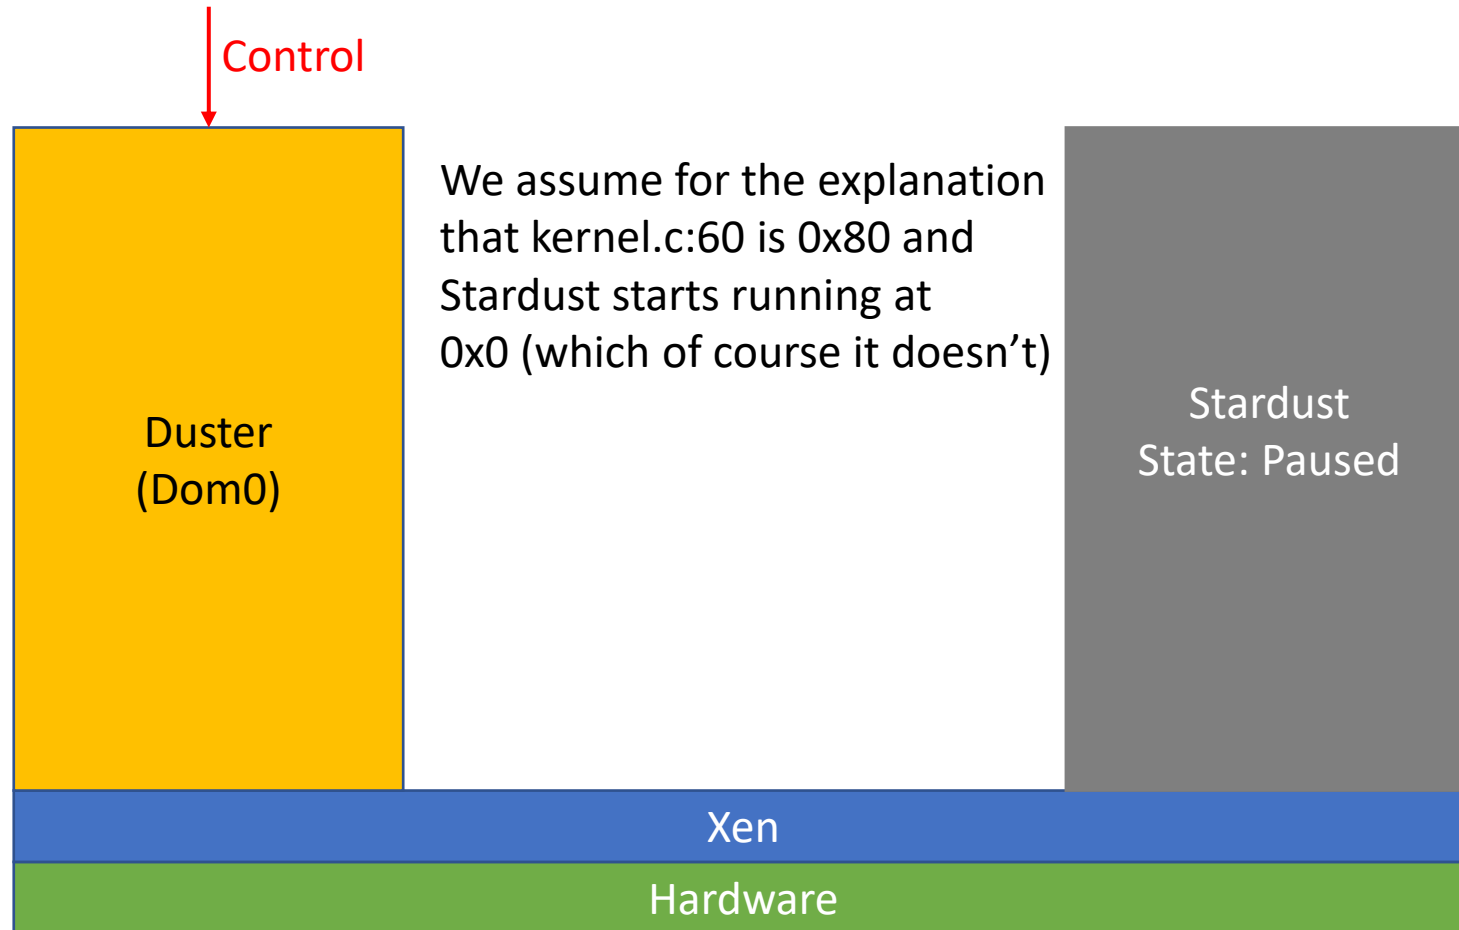
Demo Code (kernel.c)

```
59 void demo() {
60     int my_integer = 0;
61     float my_float = 0.0;
62     for (int i = 0; i < 3; i++) {
63         printf("We are on the %d iteration of the loop\n", i);
64         my_integer += 100;
65         my_float += 0.5;
66     }
67     bool my_boolean = true;
68     struct test *pointer_tester;
69     pointer_tester = malloc(sizeof(struct test));
70     pointer_tester->val = 20;
71     pointer_tester->no = 0.110;
72     pointer_tester->my_pointer = NULL;
73     printf("%d\n", pointer_tester->val);
74 }
```



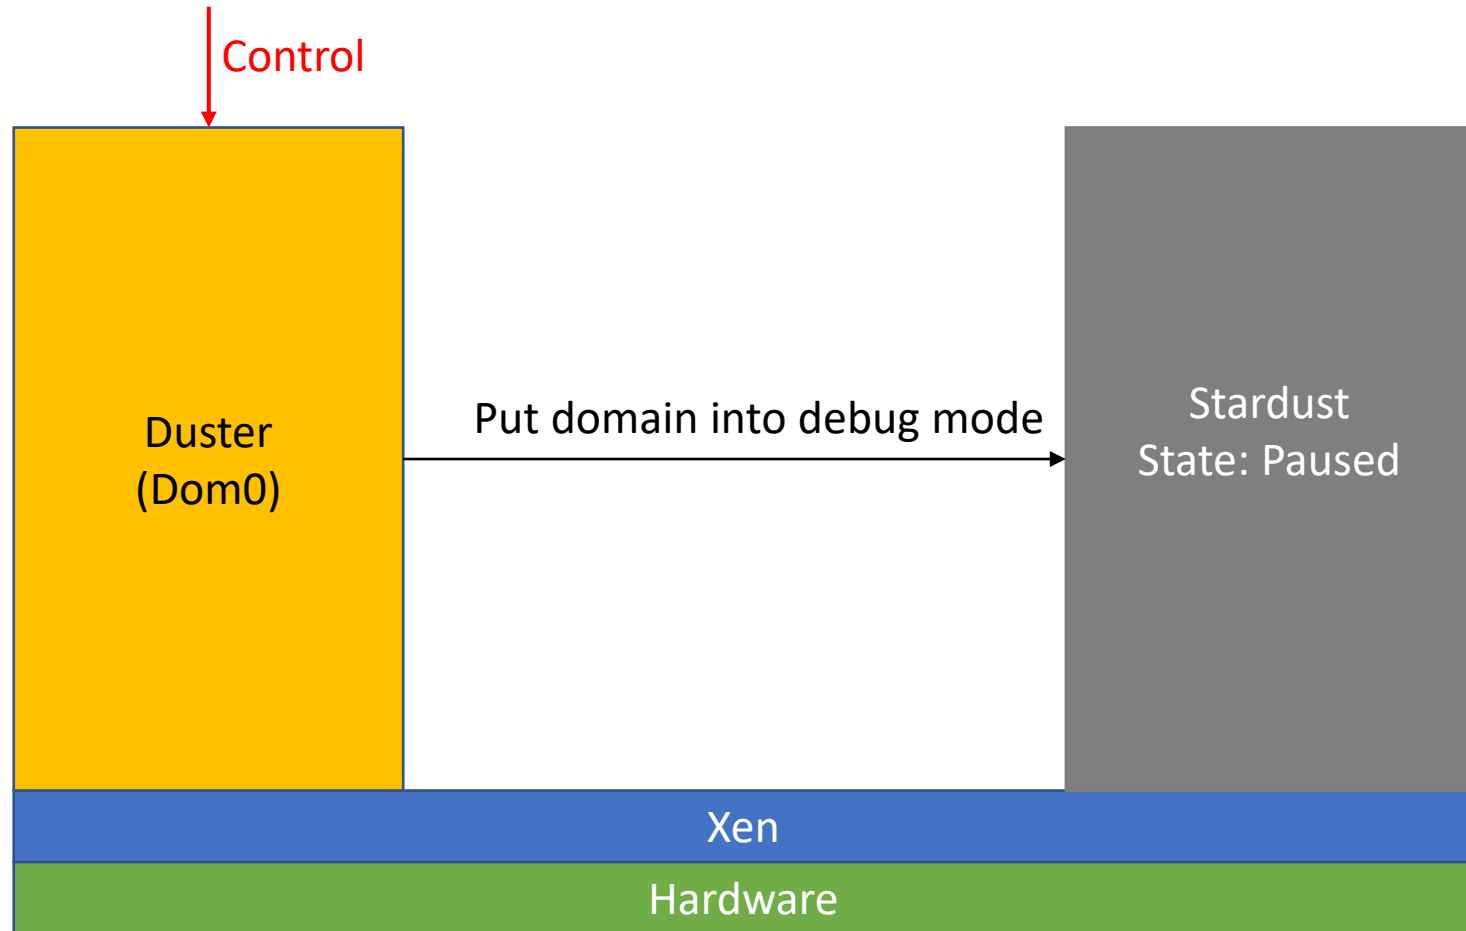


Explanation



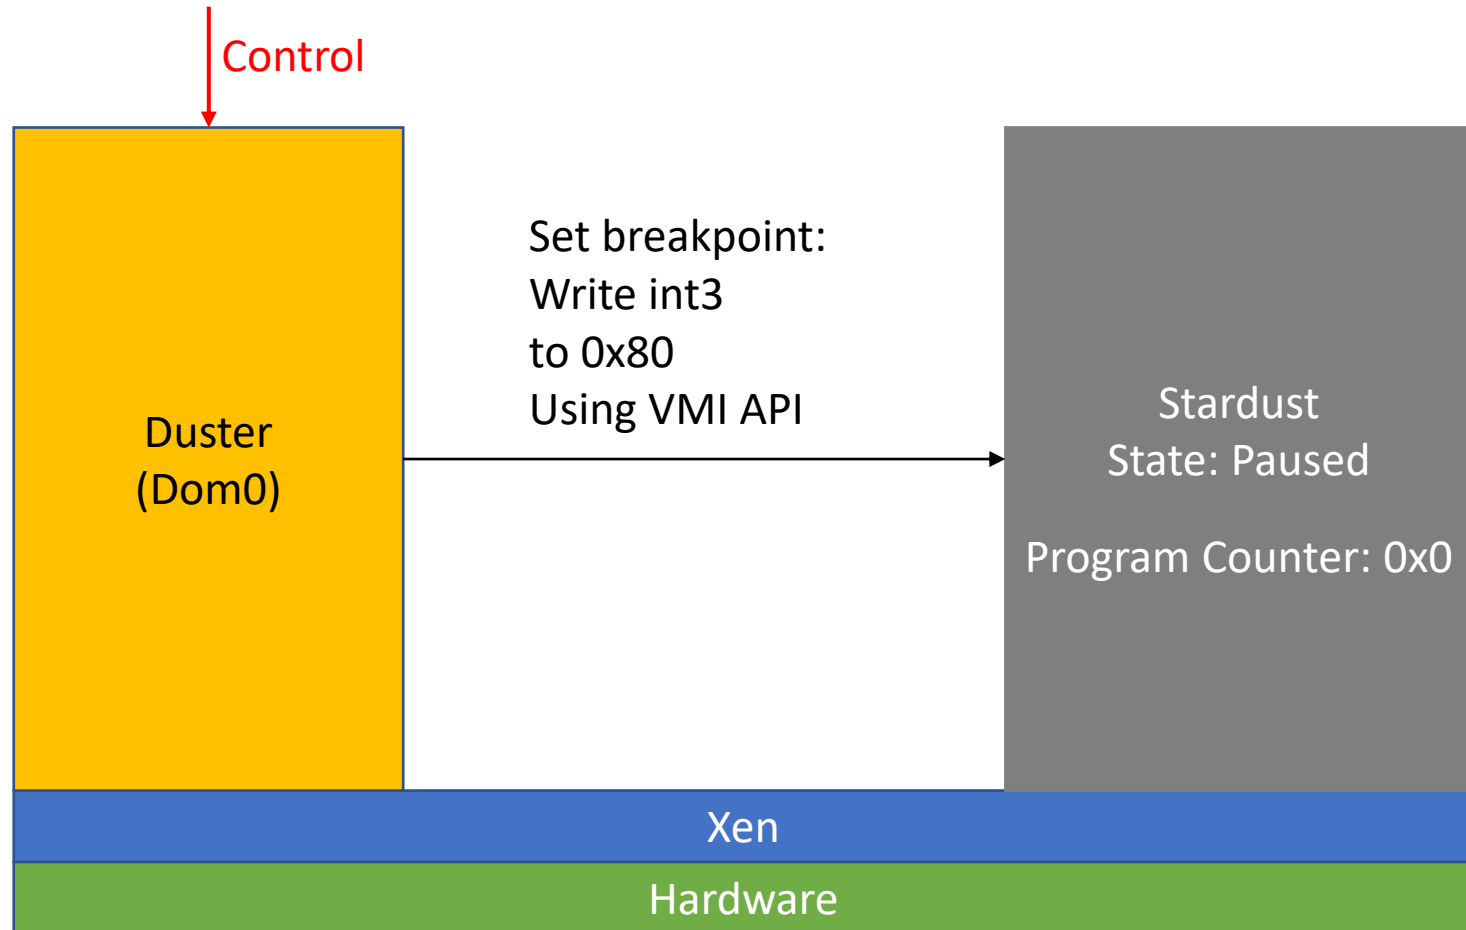


Explanation



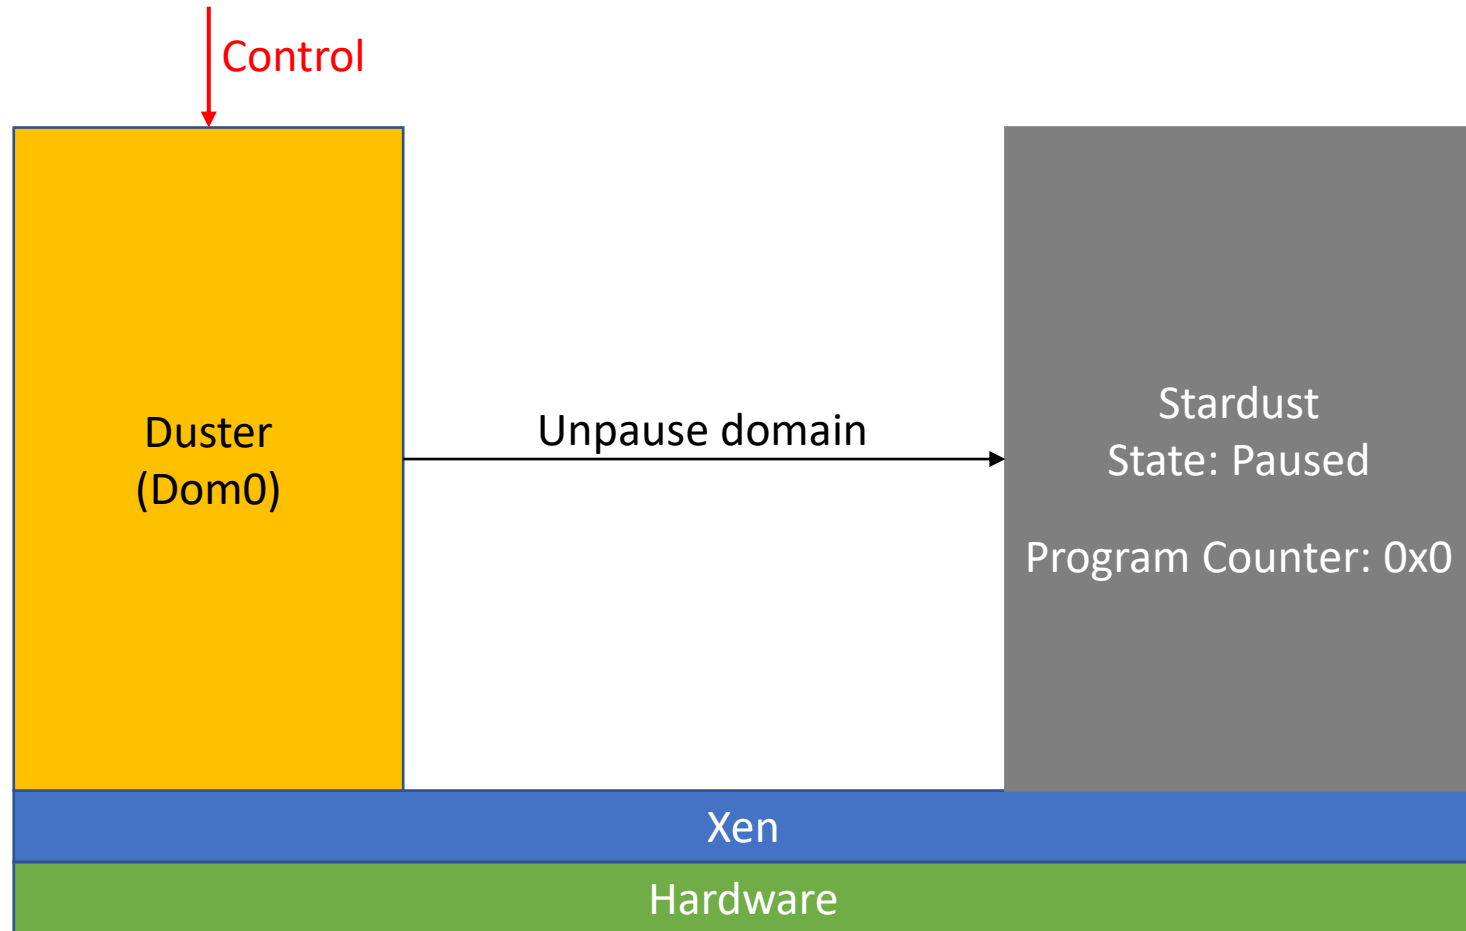


Explanation






Explanation



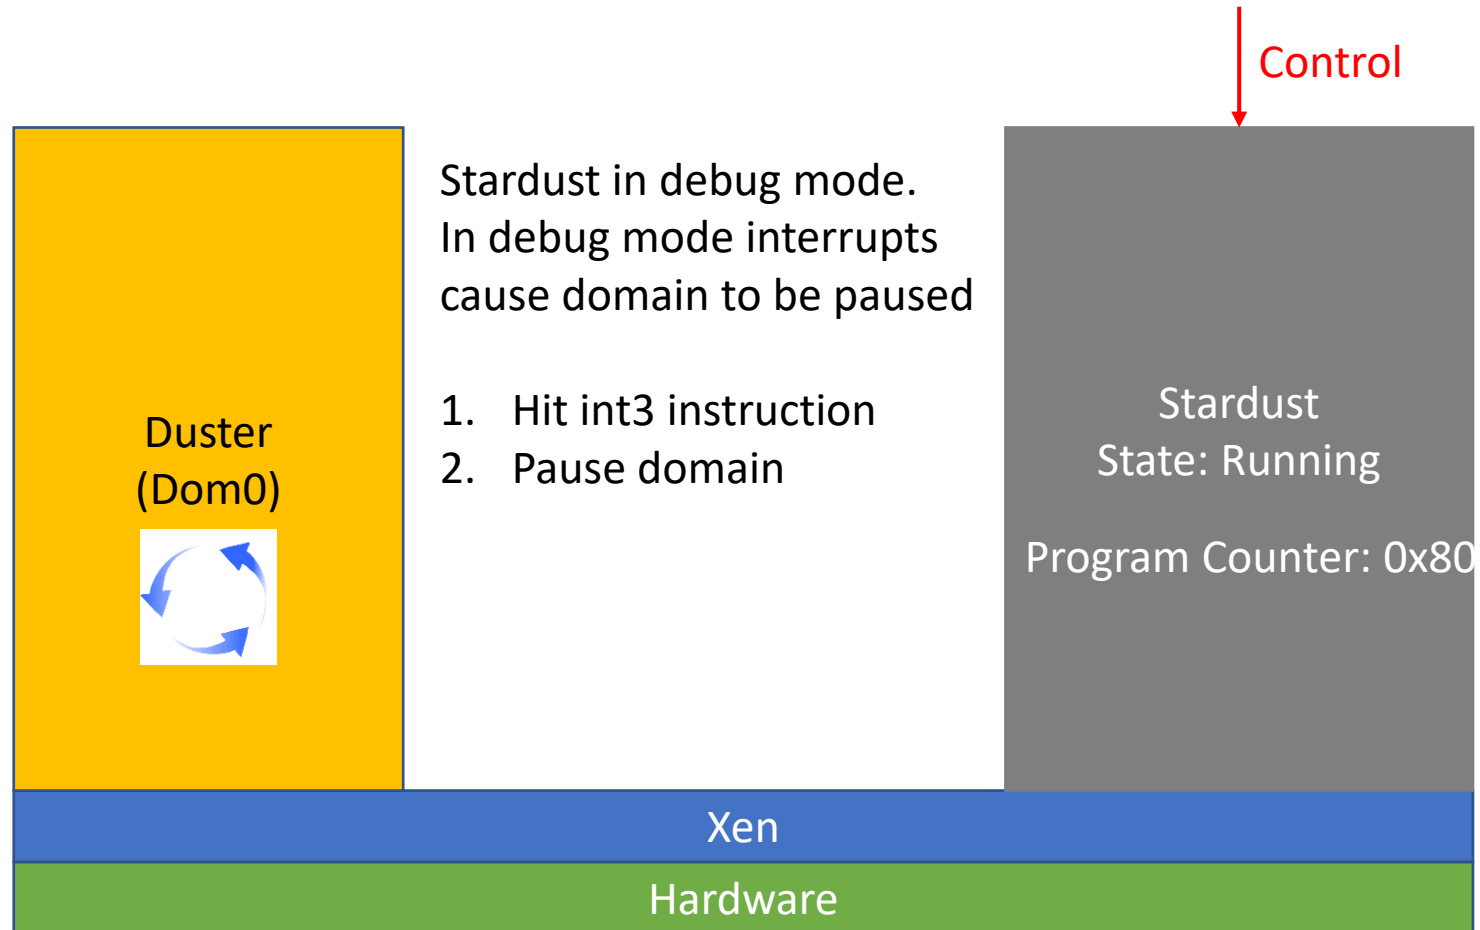


Explanation



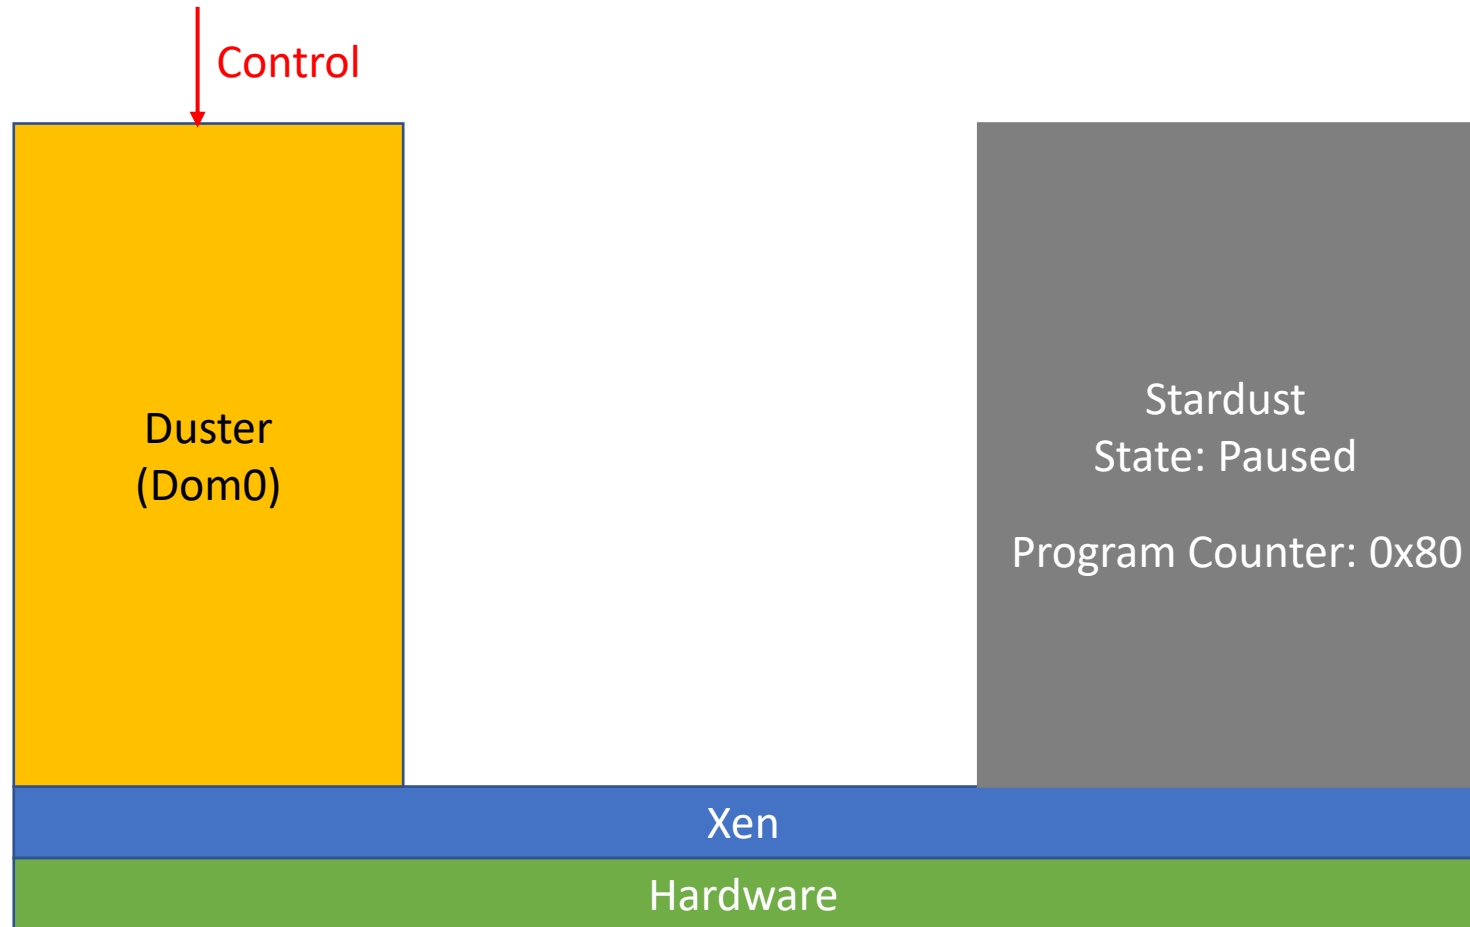


Explanation





Explanation





Demo Code

```
59 void demo() {
60     int my_integer = 0;
61     float my_float = 0.0;
62     for (int i = 0; i < 3; i++) {
63         printf("We are on the %d iteration of the loop\n", i);
64         my_integer += 100;
65         my_float += 0.5;
66     }
67     bool my_boolean = true;
68     struct test *pointer_tester;
69     pointer_tester = malloc(sizeof(struct test));
70     pointer_tester->val = 20;
71     pointer_tester->no = 0.110;
72     pointer_tester->my_pointer = NULL;
73     printf("%d\n", pointer_tester->val);
74 }
```


×

```
/home/mka/workspace/stardust/src/kernel.c:73 - printf("%d\n", pointer_tester->val);
```

break	Sets a break point at in a file (argument in the form of file.c:<line no>)
step	Steps forward one line (note a breakpoint must be set before hand)
continue	Continue to the next breakpoint
quit	Exit the debugger
read	Read a variable
der	Deference a variable



Conclusions

- We have extended the Xendbg code to support the debugging of high level (C) programming language code x86-64 Para-virtualised unikernels
- Without this (we at least) had no debug support
- Does not support some GDB operations including writing to memory, stack frame analysis
- But can:
 - Set and remove breakpoints on the source level
 - Step through code line at a time
 - Read memory using symbolic names
 - Pretty print memory based on C types



Links

- Xendbg: <https://github.com/SpencerMichaels/xendbg>
- Duster: <https://github.com/StardustOS/duster>
- Stardust: <https://github.com/StardustOS>
- Stardust
(docs): <https://stardustos.gitbook.io/docs/>